

# Vector Embeddings

---

CMSC 473/673 - NATURAL LANGUAGE PROCESSING

*Slides modified from Dr. Frank Ferraro*

# Learning Objectives

---

Correct common misconceptions about machine learning

Define a language model

Understand the use & creation of dense vector embeddings

Calculate the distance between vector embeddings

Recognize popular vector embeddings

# Misconceptions

---

Continual/Lifelong Learning vs “Regular” Machine Learning

Baselines

Determining a goal vs evaluation metrics



Language Models

# Continual Learning vs Machine Learning


## “STATIC” MACHINE LEARNING

## CONTINUAL MACHINE LEARNING

### 1) Train




### 2) Test/Deploy



# Determining how good a model is


---

## 2) Test




# Determining how good a model is: Baselines

---




# Determining how good a model is: Baselines

---



# Determining how good a model is: Evaluation Metric vs Goal

---



## What are you evaluating?

- How good is the model at translating from Mandarin to Twi?
- How accurate is the model at translating the word “potato” across languages?
- [Your questions here]



# Bonus Misconception: Data References

---

If it's cited in a paper:

## In Text

In this paper, we use ROC Stories (Mostafazadeh et al., 2016), which is a dataset...

## Reference

Mostafazadeh, N., Chambers, N., He, X., Parikh, D., Batra, D., Vanderwende, L., Kohli, P., & Allen, J. (2016). A Corpus and Cloze Evaluation for Deeper Understanding of Commonsense Stories. *Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL)*, 839–849.  
<http://www.aclweb.org/anthology/N16-1098>

# Bonus Misconception: Data References

---

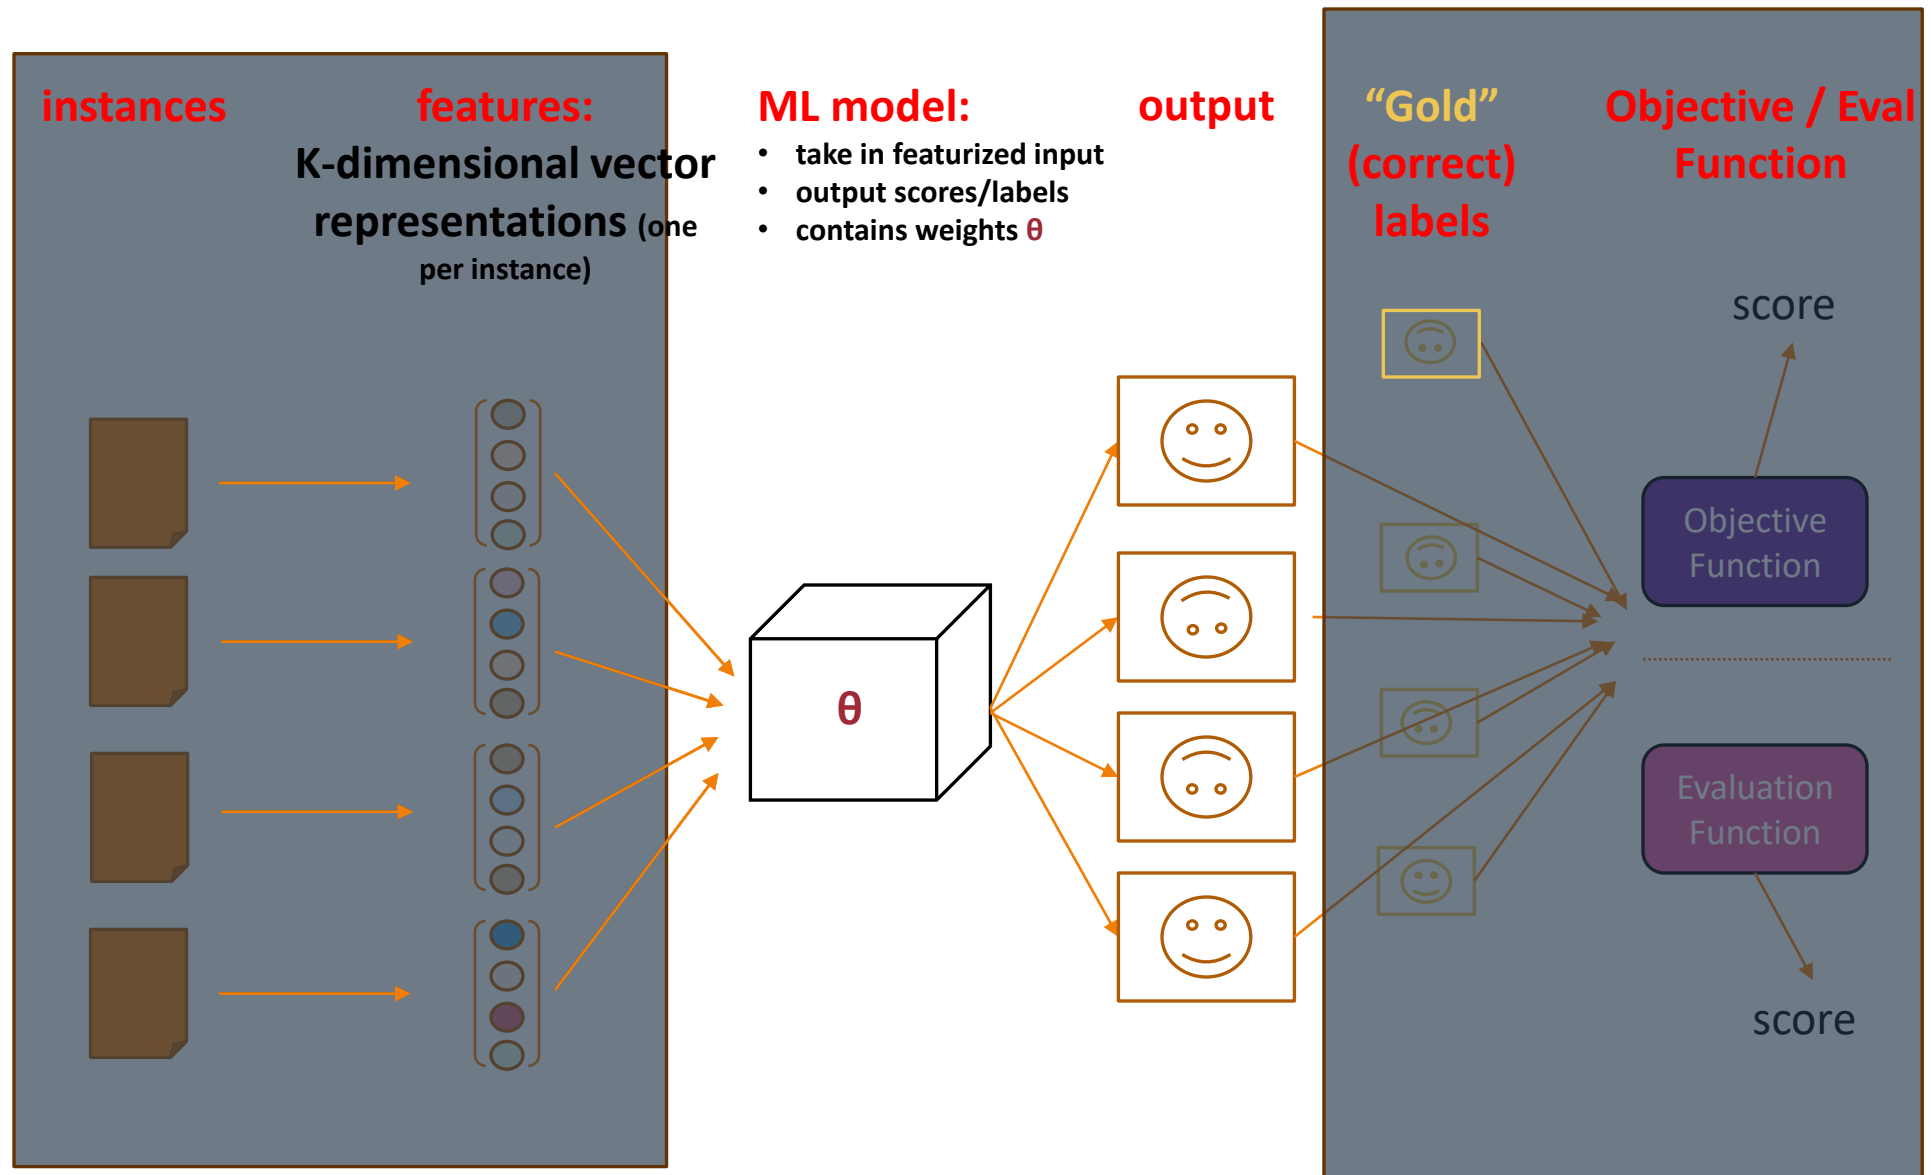
If it's not cited in a paper (i.e., just online/on Github/on 🤗):

## **In Text**

We scraped story plots from Fandom wikis<sup>1</sup>

## **Footnote**

<sup>1</sup> <https://www.fandom.com/>



# Modeling

---

Classification

$$P(y | x)$$

Language  
Model (LM)

$$P(w_t | w_{t-1}, w_{t-2} \dots)$$

A language model is used to **generate** the next word(s) given a history of words.

# Knowledge Check

---

When poll is active respond at

[PollEv.com/laramartin527](https://PollEv.com/laramartin527)

or

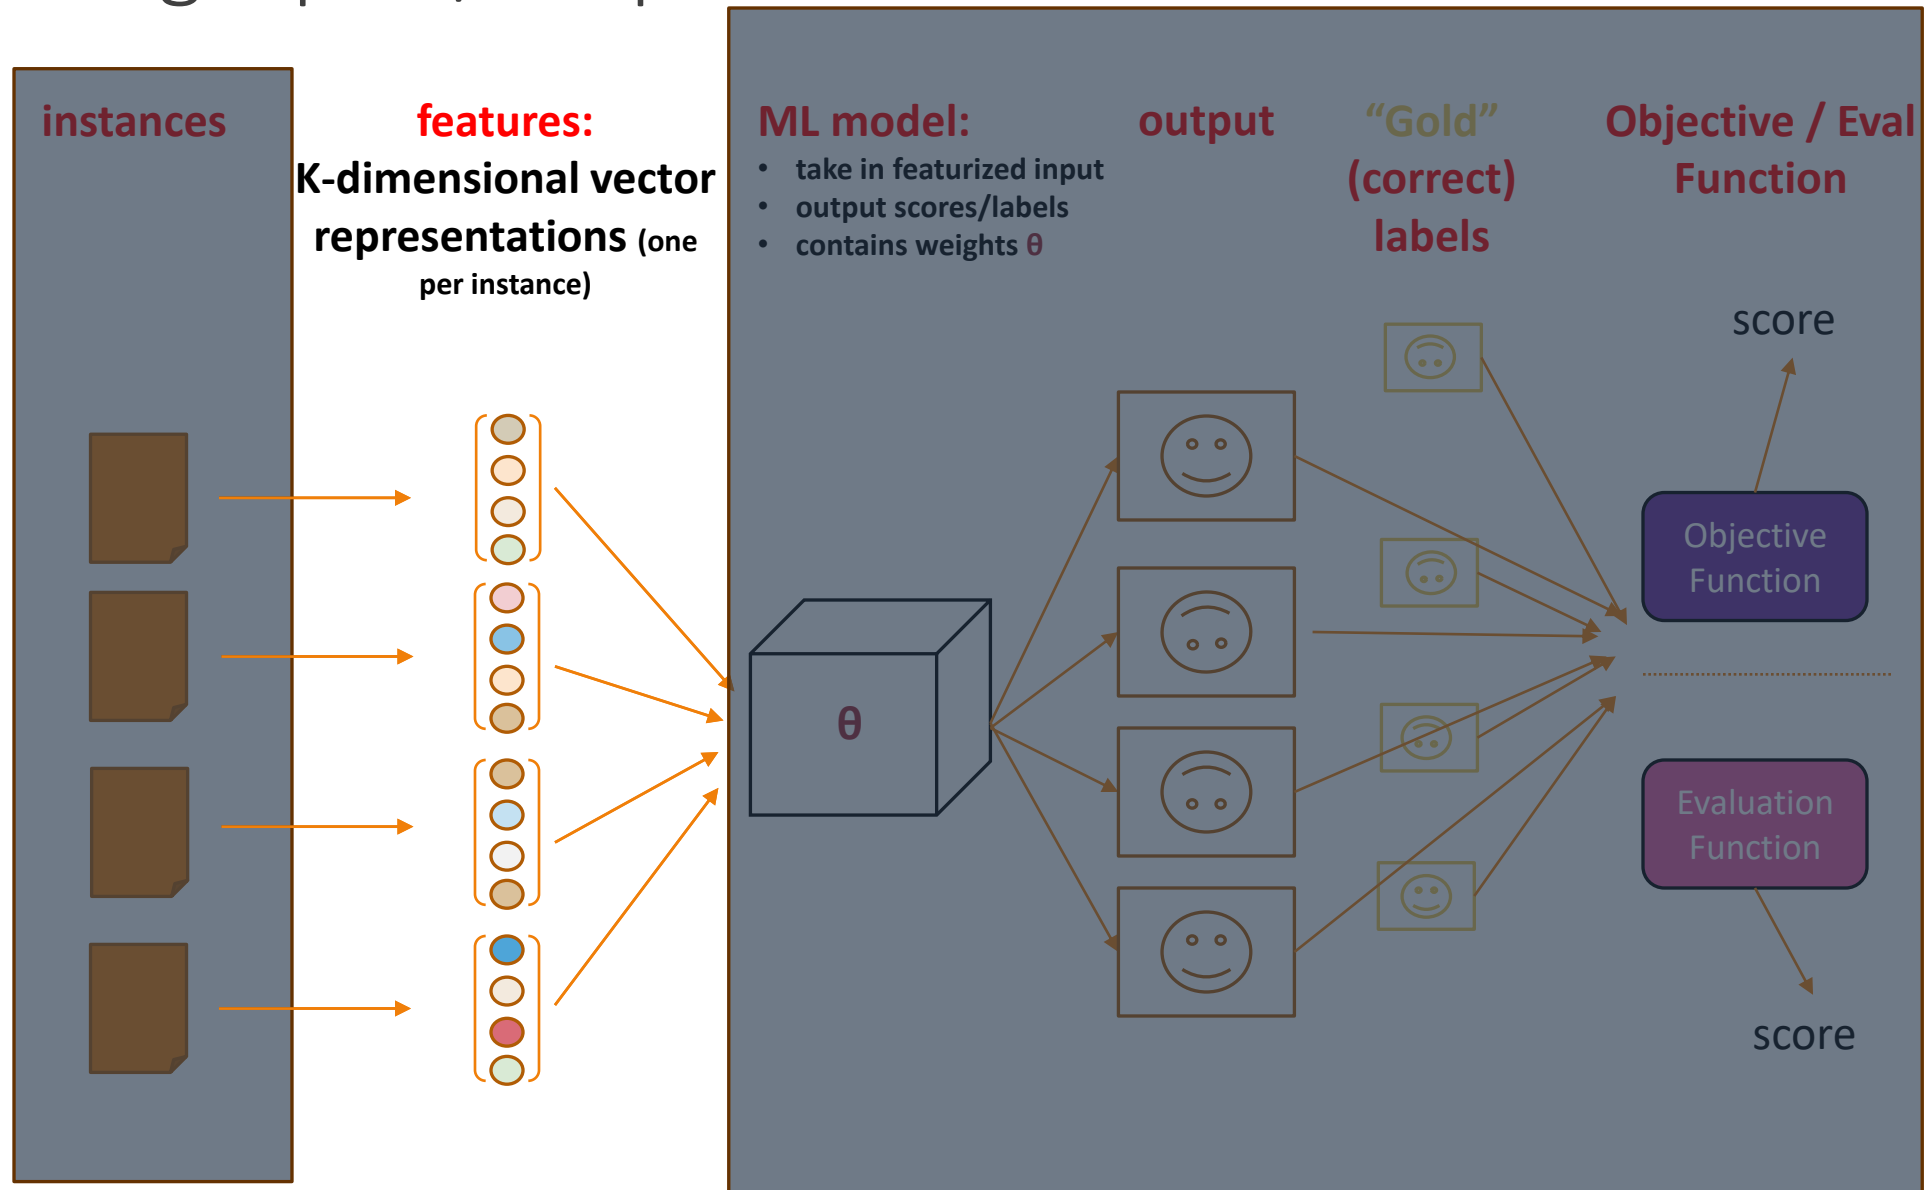
Send laramartin527 and your message to 22333



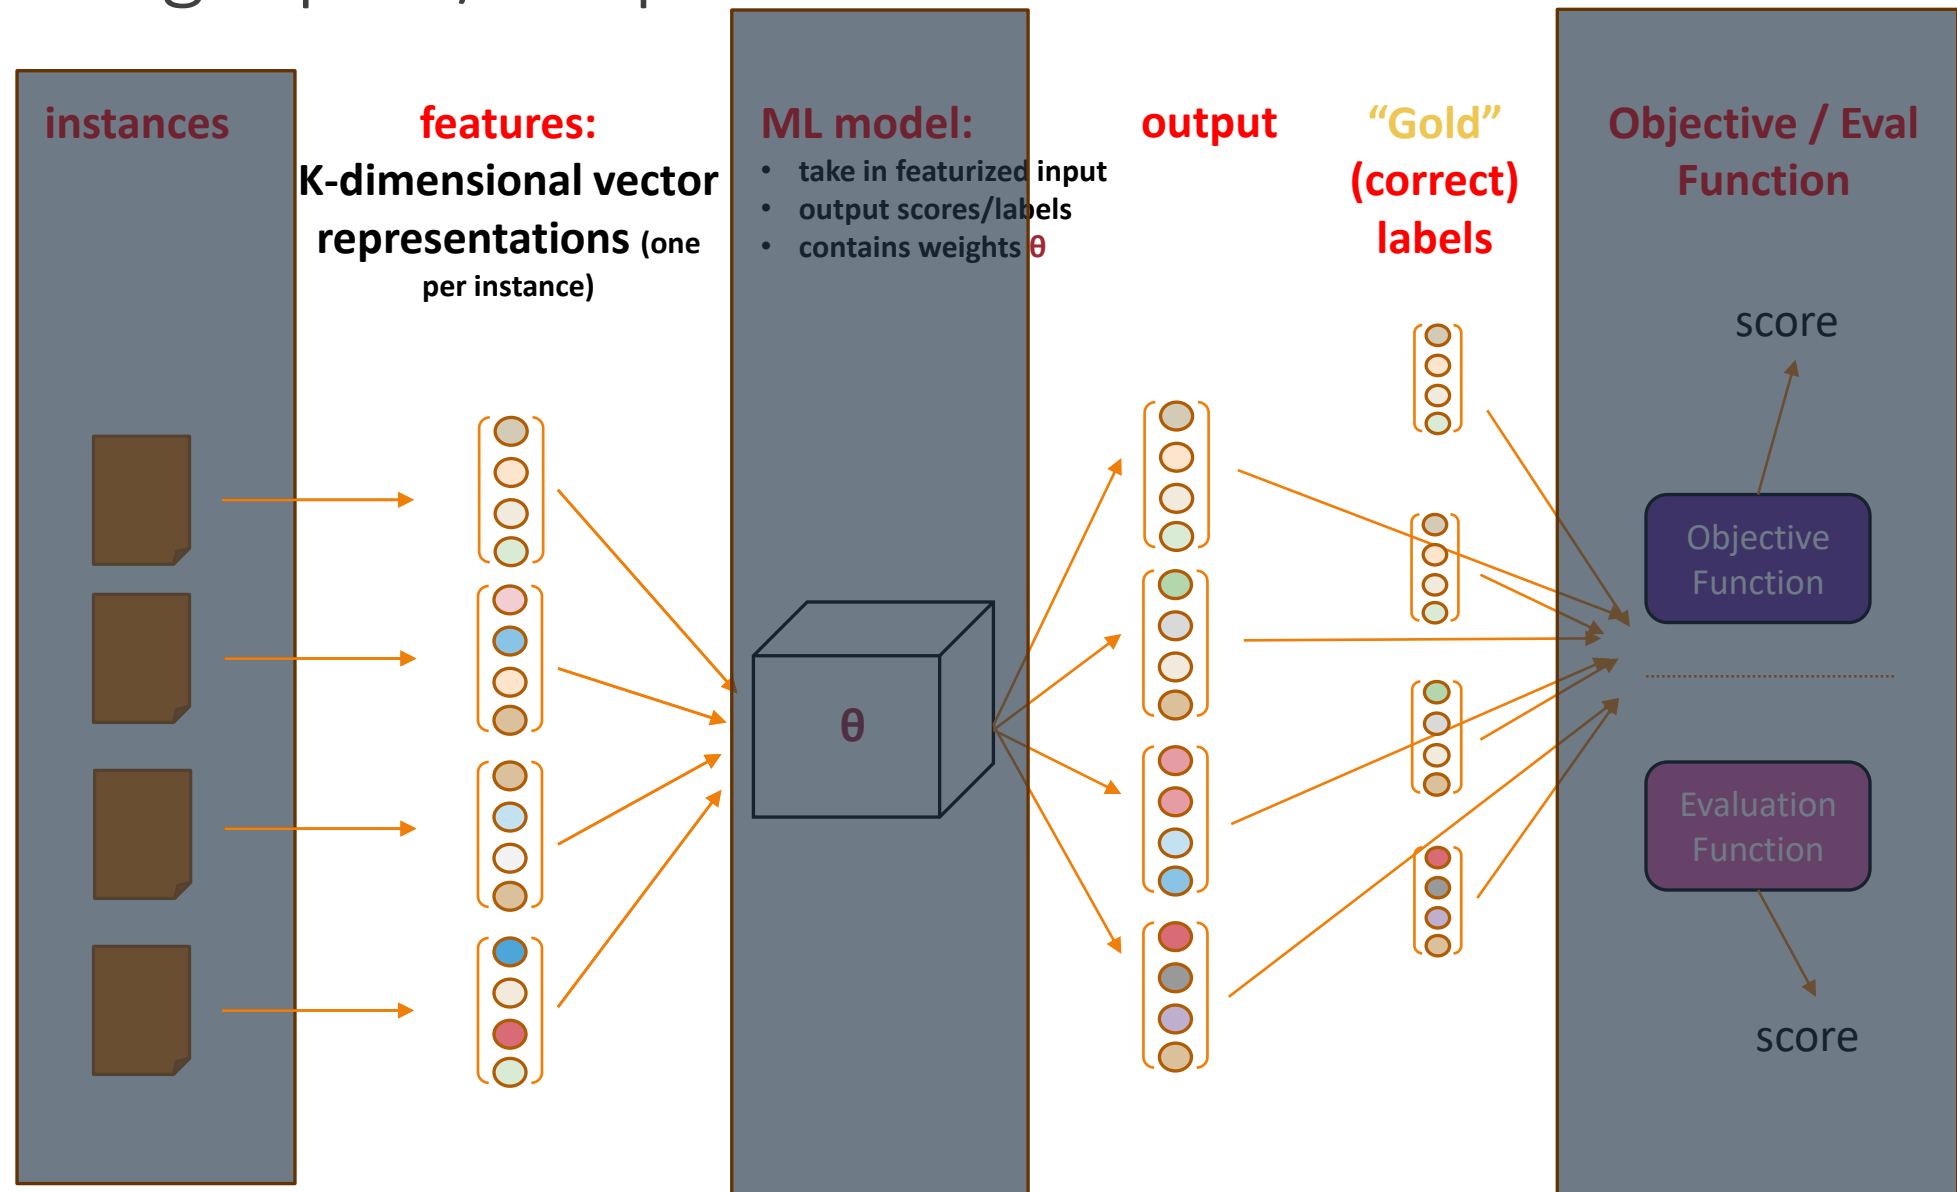
# Embeddings

---

# Representing Inputs/Outputs



# Representing Inputs/Outputs





# How have we represented words?

---

Each word is a distinct item

- Bijection between the strings and unique integer ids:
- "cat" --> 3, "kitten" --> 792 "dog" --> 17394
- Are "cat" and "kitten" similar?

Equivalently: "One-hot" encoding

- Represent each word type  $w$  with a vector the size of the vocabulary
- This vector has  $V-1$  zero entries, and 1 non-zero (one) entry

# One-Hot Encoding Example

Let our vocab be {a, cat, saw, mouse, happy}

$V = \# \text{ types} = 5$

Assign:

a	4
cat	2
saw	3
mouse	0
happy	1

How do we  
represent "cat?"

$$e_{\text{cat}} = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{pmatrix}$$

How do we  
represent  
"happy?"

$$e_{\text{happy}} = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

# The Fragility of One-Hot Encodings

## *Case Study: Maxent Plagiarism Detector*

Given two documents  $x_1$ ,  $x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

What is/are the:

Method/steps for predicting?

General formulation?

Features?



There's no way you'll  
catch me!

# Case Study: Maxent Plagiarism Detector (Feature Example)

Given two documents  $x_1, x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

Intuition: documents are more likely to be plagiarized if they have words in common

$$f_{\text{any-common-word, Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$$



# Case Study: Maxent Plagiarism Detector (Feature Example)

Given two documents  $x_1, x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

Intuition: documents are more likely to be plagiarized if they have words in common

$$f_{\text{any-common-word,Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{ngram } Z \rangle, \text{Plag.}}(x_1, x_2) = ???$$

Given two documents  $x_1, x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

Intuition: documents are more likely to be plagiarized if they have words in common

$$f_{\text{any-common-word,Plag.}}(x_1, x_2)$$

$$f_{\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ?$$

$$f_{\langle \text{ngram } Z \rangle, \text{Plag.}}(x_1, x_2) = ?$$

No problem, I'll just  
change some words!

# Case Study: Maxent Plagiarism Detector (Feature Example)

Given two documents  $x_1, x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

Intuition: documents are more likely to be plagiarized if they have words in common

$$f_{\text{any-common-word}, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{ngram } Z \rangle, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\text{synonym-of-}\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$$



Okay... but there are too many possible synonym n-grams!

# Case Study: Maxent Plagiarism Detector (Feature Example)

Given two documents  $x_1, x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

Intuition: documents are more likely to be plagiarized if they have words in common

$$f_{\text{any-common-word, Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\langle \text{ngram } Z \rangle, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\text{synonym-of-}\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$$

$$f_{\text{synonym-of-}\langle \text{ngram } Z \rangle, \text{Plag.}}(x_1, x_2) = ???$$



Hah, I win!

# Plagiarism Detection: Word Similarity?

## MAINFRAMES

Mainframes **are primarily** referred to large computers with **rapid**, advanced processing capabilities that **can execute and** perform tasks **equivalent to many** Personal Computers (PCs) machines **networked together**. It is **characterized with high quantity** Random Access Memory (RAM), very large secondary storage devices, and **high-speed** processors to cater for the needs of the computers under its service.

**Consisting of** advanced components, mainframes have the capability of running multiple large applications required by **many and** most enterprises **and organizations**. **This is** one of its advantages. Mainframes are also suitable to cater for those applications **(programs)** or files that are of very **high** demand by its users (clients). Examples of **such organizations and enterprises using mainframes are** online shopping websites **such as**

## MAINFRAMES

Mainframes **usually are** referred those computers with **fast**, advanced processing capabilities that **could perform by itself** tasks **that may require a lot of** Personal Computers (PC) Machines. **Usually mainframes would have lots of** RAMs, very large secondary storage devices, and **very fast** processors to cater for the needs of those computers under its service.

**Due to the** advanced components mainframes have, **these computers** have the capability of running multiple large applications required by most enterprises, **which is** one of its advantage. Mainframes are also suitable to cater for those applications or files that are of very **large** demand by its users (clients). Examples of these **include** the large online shopping websites **-i.e. : Ebay, Amazon, Microsoft, etc.**





# Review: Distributional Representations

---

A dense, “low”-dimensional vector representation

Many values  
are not 0 (or at  
least less  
sparse than  
one-hot)

Up till ~2013: E could be  
any size  
2013-present: E << vocab

An E-dimensional  
vector, often (but not  
always) real-valued

These are also called

- **embeddings**
- **Continuous representations**
- **(word/sentence/...) vectors**
- **Vector-space models**

# Continuous Meaning


---

The paper reflected the truth.

# Continuous Meaning

---


The paper reflected the truth.



# Continuous Meaning


---

The paper reflected the truth.



# Continuous Meaning


The paper reflected the truth.



# Continuous Meaning

---

The paper reflected the truth.





# (Some) Properties of Embeddings

Capture “like” (similar) words

target:	Redmond	Havel	ninjutsu	graffiti	capitulate
	Redmond Wash.	Vaclav Havel	ninja	spray paint	capitulation
	Redmond Washington	president Vaclav Havel	martial arts	grafitti	capitulated
	Microsoft	Velvet Revolution	swordsmanship	taggers	capitulating







# (Some) Properties of Embeddings

Capture “like” (similar) words

target:	Redmond	Havel	ninjutsu	graffiti	capitulate
	Redmond Wash.	Vaclav Havel	ninja	spray paint	capitulation
	Redmond Washington	president Vaclav Havel	martial arts	grafitti	capitulated
	Microsoft	Velvet Revolution	swordsmanship	taggers	capitulating

## Capture relationships



$$\text{vector('king')} - \text{vector('man')} + \text{vector('woman')} \approx \text{vector('queen')}$$

$$\text{vector('Paris')} - \text{vector('France')} + \text{vector('Italy')} \approx \text{vector('Rome')}$$

<https://projector.tensorflow.org/>

---

# Case Study: Maxent Plagiarism Detector (Feature Example)

Given two documents  $x_1, x_2$ , predict  $y = 1$  (plagiarized) or  $y = 0$  (not plagiarized)

Intuition: documents are more likely to be plagiarized if they have words in common



$f_{\text{common-word}, \text{Plag.}}(x_1, x_2) = ???$   
 $f_{\text{word } v, \text{Plag.}}(x_1, x_2) = ???$   
 $f_{\text{ngram } Z, \text{Plag.}}(x_1, x_2) = ???$   
 $f_{\text{synonym-of-}\langle \text{word } v \rangle, \text{Plag.}}(x_1, x_2) = ???$   
 $f_{\text{synonym-of-}\langle \text{ngram } Z \rangle, \text{Plag.}}(x_1, x_2) =$   
 $\text{get\_similarity\_with\_embeddings()}$

# Vector Representations

---

# Key Ideas

---

Vector embeddings can be used for phrases, paragraphs, or even whole documents!

1. Acquire basic contextual statistics (often counts) for each word type  $v$
2. Extract a real-valued vector  $e_v$  for each word  $v$  from those statistics

[0.00315225, 0.00315225, 0.00547597, 0.00741556, 0.00912817, 0.01068435, 0.01212381, 0.01347162, 0.01474487, 0.0159558 ]

3. Use the vectors to represent each word in later tasks

Vector from <https://www.tensorflow.org/text/tutorials/word2vec>

# Evaluating Vector Embeddings

---

# Evaluating Similarity

---

Extrinsic (task-based, end-to-end) Evaluation:

- Question Answering
- Spell Checking
- Essay grading

# Common Evaluation: Correlation between similarity ratings

---

Input: list of  $N$  word pairs  $\{(x_1, y_1), \dots, (x_N, y_N)\}$

- Each word pair  $(x_i, y_i)$  has a human-provided similarity score  $h_i$

Use your embeddings to compute an embedding similarity score  $s_i = \text{sim}(x_i, y_i)$

Compute the correlation between human and computed similarities

$$\rho = \text{Corr}((h_1, \dots, h_N), (s_1, \dots, s_N))$$

Wordsim353: 353 noun pairs rated 0-10



# Cosine: Measuring Similarity

---

Given 2 target words  $v$  and  $w$  how similar are their vectors?

Dot product or inner product from linear algebra

$$\text{dot-product}(\vec{v}, \vec{w}) = \vec{v} \cdot \vec{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

- High when two vectors have large values in same dimensions, low for orthogonal vectors with zeros in complementary distribution

Correct for high magnitude vectors

$$\frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|}$$

# Cosine Similarity

---

Divide the dot product by the length of the two vectors

$$\frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|}$$

This is the cosine of the angle between them

$$\begin{aligned}\vec{a} \cdot \vec{b} &= |\vec{a}| |\vec{b}| \cos \theta \\ \frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|} &= \cos \theta\end{aligned}$$

# Cosine Similarity


Divide the dot product by the length of the two vectors

$$\frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|}$$

This is the cosine of the angle between them

$$\vec{a} \cdot \vec{b} = |\vec{a}| |\vec{b}| \cos \theta$$

$$\frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|} = \cos \theta$$



<https://upload.wikimedia.org/wikipedia/commons/2/23/CosineSimilarity.png>

# Example: Word Similarity

$$\cos(x, y) = \frac{\sum_i x_i y_i}{\sqrt{\sum_i x_i^2} \sqrt{\sum_i y_i^2}}$$

	Dim. 1	Dim. 2	Dim. 3
apricot	2	0	0
digital	0	1	2
information	1	6	1


$$\text{cosine}(\text{apricot}, \text{information}) = \frac{2 + 0 + 0}{\sqrt{4 + 0 + 0} \sqrt{1 + 36 + 1}} = 0.1622$$

$$\text{cosine}(\text{digital}, \text{information}) = \frac{0 + 6 + 2}{\sqrt{0 + 1 + 4} \sqrt{1 + 36 + 1}} = 0.5804$$


$$\text{cosine}(\text{apricot}, \text{digital}) = \frac{0 + 0 + 0}{\sqrt{4 + 0 + 0} \sqrt{0 + 1 + 4}} = 0.0$$

# Cosine Similarity Range


---



- Angle  $\theta$  close to 0
- $\cos(\theta)$  close to 1
- **Similar vectors**



- Angle  $\theta$  close to 90
- $\cos(\theta)$  close to 0
- **Orthogonal vectors**



- Angle  $\theta$  close to 180
- $\cos(\theta)$  close to -1
- **Opposite vectors**

<https://www.learnatasci.com/glossary/cosine-similarity/>

# Other Similarity Measures

$$\text{sim}_{\text{cosine}}(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} = \frac{\sum_{i=1}^N v_i \times w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

$$\text{sim}_{\text{Jaccard}}(\vec{v}, \vec{w}) = \frac{\sum_{i=1}^N \min(v_i, w_i)}{\sum_{i=1}^N \max(v_i, w_i)}$$

$$\text{sim}_{\text{Dice}}(\vec{v}, \vec{w}) = \frac{2 \times \sum_{i=1}^N \min(v_i, w_i)}{\sum_{i=1}^N (v_i + w_i)}$$

$$\text{sim}_{\text{JS}}(\vec{v} || \vec{w}) = D(\vec{v} | \frac{\vec{v} + \vec{w}}{2}) + D(\vec{w} | \frac{\vec{v} + \vec{w}}{2})$$

# Adding Morphology, Syntax, and Semantics to Embeddings

---

- Lin (1998): “Automatic Retrieval and Clustering of Similar Words”
- Padó and Lapata (2007): “Dependency-based Construction of Semantic Space Models”
- Levy and Goldberg (2014): “Dependency-Based Word Embeddings”
- Cotterell and Schütze (2015): “Morphological Word Embeddings”
- Ferraro et al. (2017): “Frame-Based Continuous Lexical Semantics through Exponential Family Tensor Factorization and Semantic Proto-Roles”
- and many more...

# Common Continuous Representations

---



# Shared Intuition

---

Model the meaning of a word by “embedding” in a vector space

The meaning of a word is a vector of numbers

Contrast: word meaning is represented in many computational linguistic applications by a vocabulary index (“word number 545”) or the string itself

# Three Common Kinds of Embedding Models

---

1. Co-occurrence matrices
2. Matrix Factorization: Singular value decomposition/Latent Semantic Analysis, Topic Models
3. Neural-network-inspired models (skip-grams, CBOW)

# Three Common Kinds of Embedding Models

---


1. Co-occurrence matrices
2. Matrix Factorization: Singular value decomposition/Latent Semantic Analysis, Topic Models
3. Neural-network-inspired models (skip-grams, CBOW)

Co-occurrence matrices can be used in their own right, but they're most often used as inputs (directly or indirectly) to the matrix factorization or neural approaches

# Co-occurrence Matrix

---

Acquire basic contextual statistics  
(often counts) for each word type  $v$  via  
*correlate*.




# Co-occurrence Matrix

Acquire basic contextual statistics  
(often counts) for each word type  $v$  via  
*correlate*:

For example:

documents

- Record how often a word occurs in each document



# correlates =  
# documents

# Co-occurrence Matrix


Acquire basic contextual statistics  
(often counts) for each word type  $v$  via  
*correlate*:

For example:

documents

surrounding context words

- Record how often  $v$  occurs with other word types  $u$



# correlates =  
# word types

# Co-occurrence Matrix

Acquire basic contextual statistics  
(often counts) for each word type  $v$  via  
*correlate*:

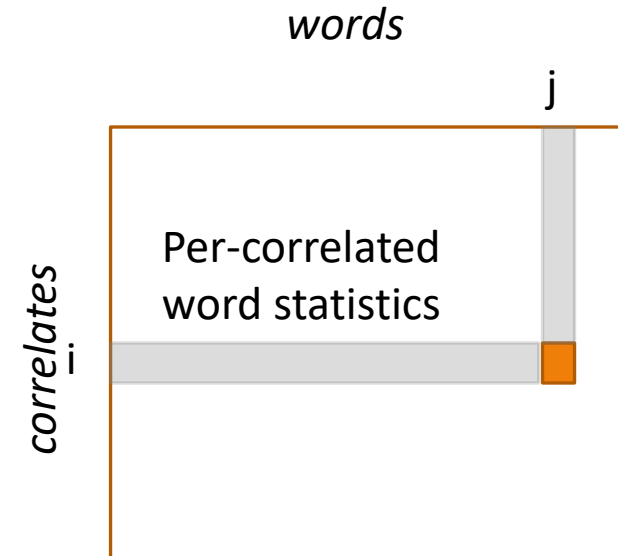
For example:

documents

surrounding context words

linguistic annotations (POS tags,  
syntax)

...



*Assumption: Two words  
are similar if their  
vectors are similar*

# “Acquire basic contextual statistics (often counts) for each word type $v$ ”

---

Two basic, initial counting approaches

- Record which words appear in which documents
- Record which words appear together

These are good first attempts, but with some large downsides



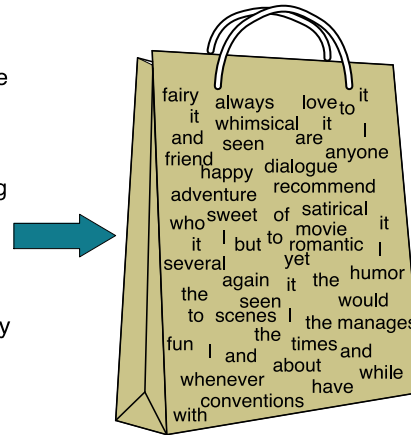
# “You shall know a word by the company it keeps!” Firth (1957)

document (↓)-word (→) count matrix

	battle	soldier	fool	clown
<i>As You Like It</i>	1	2	37	6
<i>Twelfth Night</i>	1	2	58	117
<i>Julius Caesar</i>	8	12	1	0
<i>Henry V</i>	15	36	5	0

basic bag-of-  
words  
counting

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



it 6  
I 5  
the 4  
to 3  
and 3  
seen 2  
yet 1  
would 1  
whimsical 1  
times 1  
sweet 1  
satirical 1  
adventure 1  
genre 1  
fairy 1  
humor 1  
have 1  
great 1  
... ..

# “You shall know a word by the company it keeps!” Firth (1957)

---

document (↓)-word (→) count matrix

	battle	soldier	fool	clown
<i>As You Like It</i>	1	2	37	6
<i>Twelfth Night</i>	1	2	58	117
<i>Julius Caesar</i>	8	12	1	0
<i>Henry V</i>	15	36	5	0

*Assumption: Two documents are similar if their vectors are similar*

# “You shall know a word by the company it keeps!” Firth (1957)

---

document (↓)-word (→) count matrix

	battle	soldier	fool	clown
<i>As You Like It</i>	1	2	37	6
<i>Twelfth Night</i>	1	2	58	117
<i>Julius Caesar</i>	8	12	1	0
<i>Henry V</i>	15	36	5	0

*Assumption: Two words are similar if their vectors are similar???*

# “You shall know a word by the company it keeps!” Firth (1957)

---

document (↓)-word (→) count matrix

	battle	soldier	fool	clown
<i>As You Like It</i>	1	2	37	6
<i>Twelfth Night</i>	1	2	58	117
<i>Julius Caesar</i>	8	12	1	0
<i>Henry V</i>	15	36	5	0

*Assumption: Two words are similar if their vectors are similar*

*Issue: Count word vectors are very large, sparse, and skewed!*

# “You shall know a word by the company it keeps!” Firth (1957)

---

**context** (↓)-**word** (→) count matrix

	apricot	pineapple	digital	information
aardvark	0	0	0	0
computer	0	0	2	1
data	0	10	1	6
pinch	1	1	0	0
result	0	0	1	4
sugar	1	1	0	0

*Context: those other words within a small “window” of a target word*

# “You shall know a word by the company it keeps!” Firth (1957)

**context** (↓)-**word** (→) count matrix

	apricot	pineapple	digital	information
aardvark	0	0	0	0
computer	0	0	2	1
data	0	10	1	6
pinch	1	1	0	0
result	0	0	1	4
sugar	1	1	0	0

*Context: those other words within a small “window” of a target word*

a cloud **[** computer stores digital data on **]** a remote computer

# “You shall know a word by the company it keeps!” Firth (1957)

---

**context** (↓)-**word** (→) count matrix

	apricot	pineapple	digital	information
aardvark	0	0	0	0
computer	0	0	2	1
data	0	10	1	6
pinch	1	1	0	0
result	0	0	1	4
sugar	1	1	0	0

The size of windows depends on your goals

The shorter the windows , the more **syntactic** the representation

± 1-3 more “syntax-y”

The longer the windows, the more **semantic** the representation

± 4-10 more “semantic-y”

# “You shall know a word by the company it keeps!” Firth (1957)

---

**context** (↓)-**word** (→) count matrix

	apricot	pineapple	digital	information
aardvark	0	0	0	0
computer	0	0	2	1
data	0	10	1	6
pinch	1	1	0	0
result	0	0	1	4
sugar	1	1	0	0

*Context: those other words within a small “window” of a target word*

*Assumption: Two words are similar if their vectors are similar*

*Issue: Count word vectors are very large, sparse, and skewed!*



# Pointwise Mutual Information (PMI): Dealing with Problems of Raw Counts

---

Raw word frequency is not a great measure of association between words

It's very skewed: "the" and "of" are very frequent, but maybe not the most discriminative

We'd rather have a measure that asks whether a context word is **particularly informative** about the target word.

(Positive) Pointwise Mutual Information ((P)PMI)

## Pointwise mutual information:

Do events  $x$  and  $y$  co-occur more than if they were independent?

probability words  $x$  and  $y$  occur together  
(in the same context/window)

$$\text{PMI}(x, y) = \log \frac{p(x, y)}{p(x)p(y)}$$

probability that  
word  $x$  occurs

probability that  
word  $y$  occurs

# Advanced: Equivalent PMI Computations

---

Intuition: Do words  $x$  and  $y$  co-occur more than if they were independent?

$$\text{PMI}(x, y) = \log \frac{p(x, y)}{p(x)p(y)} = \log \frac{p(y | x)}{p(y)} = \log \frac{p(x | y)}{p(x)}$$

# “Noun Classification from Predicate-Argument Structure,” Hindle (1990)

“**drink it**” is more common than “**drink wine**”

“**wine**” is a better “drinkable” thing than “**it**”

Object of “drink”	Count	PMI
<b>it</b>	3	1.3
anything	3	5.2
<b>wine</b>	2	9.3
tea	2	11.8
liquid	2	10.5

# Three Common Kinds of Embedding Models

---

Learn more in:

- Your project
- Paper (673)

1. Co-occurrence Models
  - Other classes (478/678)
2. Matrix Factorization: Singular value decomposition/Latent Semantic Analysis, Topic Models
3. Neural-network-inspired models (skip-grams, CBOW)

# Three Common Kinds of Embedding Models

---

1. Co-occurrence matrices
2. Matrix Factorization: Singular value decomposition/Latent Semantic Analysis, Topic Models
3. Neural-network-inspired models (skip-grams, CBOW)

# Word2Vec

---

Mikolov et al. (2013; NeurIPS): “Distributed Representations of Words and Phrases and their Compositionality”

Revisits the context-word approach

Learn a model  $p(c \mid w)$  to predict a context word from a target word

Learn two types of vector representations

- $h_c \in \mathbb{R}^E$ : vector embeddings for each context word
- $v_w \in \mathbb{R}^E$ : vector embeddings for each target word

$$p(c \mid w) \propto \exp(h_c^T v_w)$$

# Word2Vec

**context** (↓)-**word** (→) count matrix

	apricot	pineapple	digital	information
aardvark	0	0	0	0
computer	0	0	2	1
data	0	10	1	6
pinch	1	1	0	0
result	0	0	1	4
sugar	1	1	0	0

*Context: those other words within a small “window” of a target word*

$$\max_{h,v} \sum_{c,w \text{ pairs}} \text{count}(c, w) \log p(c | w)$$

# Word2Vec

**context** (↓)-**word** (→) count matrix

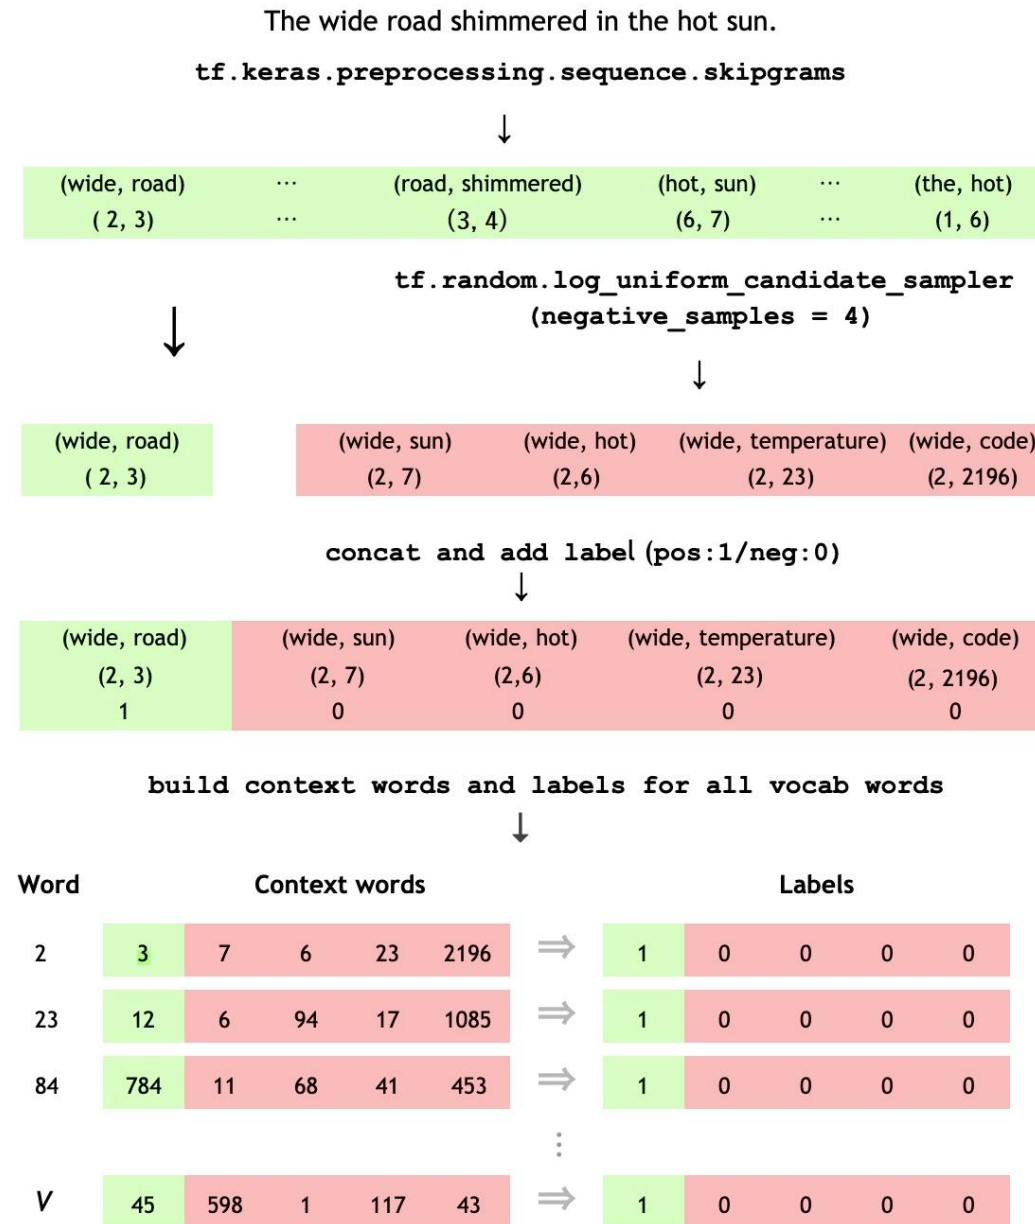
	apricot	pineapple	digital	information
aardvark	0	0	0	0
computer	0	0	2	1
data	0	10	1	6
pinch	1	1	0	0
result	0	0	1	4
sugar	1	1	0	0

*Context: those other words within a small “window” of a target word*

$$\max_{h,v} \sum_{c,w \text{ pairs}} \text{count}(c, w) \left[ h_c^T v_w - \log \left( \sum_u \exp(h_u^T v_w) \right) \right]$$



# Example (Tensorflow)



<https://www.tensorflow.org/text/tutorials/word2vec>

# Word2Vec has Inspired a Lot of Work

---

## Off-the-shelf embeddings

- <https://code.google.com/archive/p/word2vec/>

## Off-the-shelf implementations

- <https://radimrehurek.com/gensim/models/word2vec.html>

## Follow-on work

- J. Pennington, R. Socher, and C. D. Manning, “**GLoVe: Global Vectors for Word Representation**,” in *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Doha, Qatar, 2014, pp. 1532–1543. doi: [10.3115/v1/D14-1162](https://doi.org/10.3115/v1/D14-1162).
  - <https://nlp.stanford.edu/projects/glove/>
- Many others
- 15000+ citations

# FastText

---

P. Bojanowski, E. Grave, A. Joulin, and T. Mikolov, “**Enriching Word Vectors with Subword Information**,” *Transactions of the Association for Computational Linguistics*, vol. 5, pp. 135–146, 2017, doi: [10.1162/tac1\\_a\\_00051](https://doi.org/10.1162/tac1_a_00051).

Main idea: learn **character n-gram embeddings** for the target word (not context) and modify the word2vec model to use these

Pre-trained models in 150+ languages

- <https://fasttext.cc>

# FastText Details

---

Main idea: learn **character n-gram embeddings** and for the target word (not the context) modify the word2vec model to use these

Original word2vec:

$$p(c | w) \propto \exp(h_c^T v_w)$$

FastText:

$$p(c | w) \propto \exp\left(h_c^T \left(\sum_{\text{n-gram } g \text{ in } w} z_g\right)\right)$$

# FastText Details

---

Main idea: learn **character n-gram embeddings** and for the target word (not the context) modify the word2vec model to use these

$$p(c | w) \propto \exp \left( h_c^T \left( \sum_{n\text{-gram } g \text{ in } w} z_g \right) \right)$$

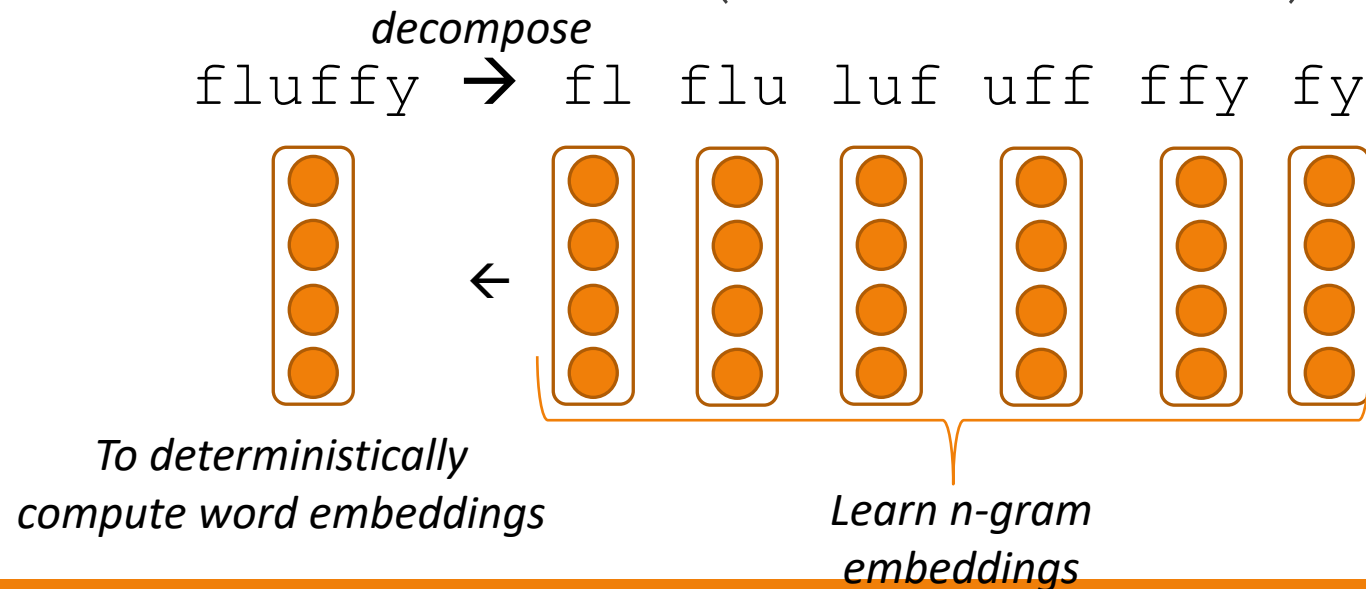
*decompose*  
fluffy  $\rightarrow$  fl flu luf uff ffy fy

Sub-word units like this have become an important part of today's NLP work!

# FastText Details

Main idea: learn **character n-gram embeddings** and for the target word (not the context) modify the word2vec model to use these

$$p(c | w) \propto \exp \left( h_c^T \left( \sum_{\text{n-gram } g \text{ in } w} z_g \right) \right)$$







# Contextual Word Embeddings

---

Word2vec-based models are not context-dependent

Single word type → single word embedding

If a single word type can have different meanings...

bank, bass, plant,...

... why should we only have one embedding?

Entire task devoted to classifying these meanings:

**Word Sense Disambiguation**



# Contextual Word Embeddings

---

Growing interest in this

Off-the-shelf is a bit more difficult

- Download and run a model
- Can't just download a file of embeddings

Two to know about (with code):

- ELMo: “Deep contextualized word representations” Peters et al. (2018; NAACL)
- <https://allennlp.org/elmo>
- BERT: “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding” Devlin et al. (2019; NAACL)
- <https://github.com/google-research/bert>

